

LINMA2710 - Scientific Computing

Graphics processing unit (GPU)

P.-A. Absil and B. Legat

Full Width Mode Present Mode

≡ Table of Contents

Introduction

Examples

Reduction on GPU

Sources

- [OpenCL.jl](#)
- [HandsOnOpenCL](#)
- [Optimizing Parallel Reduction in CUDA](#)
- [Parallel Computation Patterns \(Reduction\)](#)
- [Profiling, debugging and optimization](#)
- [How to use TAU for Performance Analysis](#)

Introduction ↵

Context ↵

Critical feedback from ArgTools module

```
RequestError: HTTP/2 429 while requesting  
https://upload.wikimedia.org/wikipedia/commons/2/25/WebGL_Logo.svg
```

Drill down into execution genealogy...


```
1 hbox([  
2     md"""  
3 * Most *dedicated* GPUs produced by  
$(img("https://upload.wikimedia.org/wikipedia/commons/a/a4/NVIDIA_logo.svg",  
:height => "15pt")) and  
$(img("https://upload.wikimedia.org/wikipedia/commons/7/7c/AMD_Logo.svg", :height  
=> "15pt"))  
4 * *Integrated* GPUs by  
$(img("https://upload.wikimedia.org/wikipedia/commons/6/6a/Intel_logo_%282020%2C_dar  
k_blue%29.svg", :height => "15pt")) used in laptops to reduce power consumption  
5 * Designed for 3D rendering through ones of the APIs :  
$(img("https://upload.wikimedia.org/wikipedia/commons/7/7f/Microsoft-DirectX-Logo-  
wordmark.svg", :height => "20pt")),  
$(img("https://upload.wikimedia.org/wikipedia/commons/2/21/OpenGL_logo.svg",  
:height => "20pt")),  
$(img("https://upload.wikimedia.org/wikipedia/commons/2/25/WebGL_Logo.svg", :height  
=> "20pt")),  
$(img("https://upload.wikimedia.org/wikipedia/commons/2/2f/WebGPU_logo.svg",  
:height => "25pt")),  
$(img("https://upload.wikimedia.org/wikipedia/commons/f/fe/Vulkan_logo.svg",  
:height => "20pt")) or Apple's Metal  
$(img("https://upload.wikimedia.org/wikipedia/commons/8/8d/Metal_3_Logo.png",  
:height => "20pt"))  
6 * Illustration on the right is from [Charge's film]  
(https://studio.blender.org/blog/charge-poster/?utm_medium=homepage), it shows how
```

```
3D modeling works.  
7 """,  
8     img("https://upload.wikimedia.org/wikipedia/commons/f/fd/Charge-  
9      movie_poster.jpg", :width => "120"),  
9 ])
```

Analyzing stealth implementations

General-Purpose computing on GPU (GPGPU) ↗

Also known as *compute shader* as they abuse the programmable shading of GPUs by treating the data as texture maps.



Critical feedback from ArgTools module

```
RequestError: HTTP/2 429 while requesting  
https://upload.wikimedia.org/wikipedia/commons/b/b9/Nvidia_CUDA_Logo.jpg
```


Drill down into execution genealogy...

```
1 grid([  
2   md"Hardware-specific"  
3   img("https://upload.wikimedia.org/wikipedia/commons/b/b9/Nvidia_CUDA_Logo.jpg",  
4     :height => "70pt")  
5   img("https://upload.wikimedia.org/wikipedia/commons/7/7b/ROCM_logo.png",  
6     :height => "60pt") md"""\n    $img("https://upload.wikimedia.org/wikipedia/commons/6/6a/Intel_logo_%282020%2C  
7     _dark_blue%29.svg", :height => "30pt"))  
8   $img("https://upload.wikimedia.org/wikipedia/en/f/fa/OneAPI-rgb-3000.png",  
9     :height => "60pt"))"""\n3   md"Common interface"  
4   img("https://upload.wikimedia.org/wikipedia/commons/4/4d/OpenCL_logo.svg")  
5   img("https://upload.wikimedia.org/wikipedia/commons/1/12/SYCL_logo.svg")  
6   img("https://d29g4g2dyqv443.cloudfront.net/sites/default/files/akamai/designwork  
7     s/blog1/OpenACC-logo.png", :height => "50pt")  
8 ])
```

Analyzing stealth implementations

Standard Portable Intermediate Representation (SPIR)

Similar to LLVM IR : Intermediate representation for accelerated computation.



Hierarchy

Critical feedback from ArgTools module

```
RequestError: HTTP/2 429 while requesting
https://upload.wikimedia.org/wikipedia/de/9/96/Platform_architecture_2009-11-08.svg
```

Drill down into execution genealogy...

```
1 hbox([
2   md"""
3   * CPUs:
4     - All CPUs part of same device
5     - 1 Compute Unit per core
```

```
6      - Number of processing elements equal to SIMD width
7 * GPUs:
8   - One device per GPU
9 """
10    img("https://upload.wikimedia.org/wikipedia/de/9/96/Platform_architecture_2009-
11-08.svg", :width => "400pt"),
11 ])
```

• Analyzing stealth implementations

compute device	compute unit	processing element
get_global_id	get_group_id	get_local_id
get_global_size	get_num_groups	get_local_size

Memory ↵

Critical feedback from ArgTools module

```
RequestError: HTTP/2 429 while requesting
https://upload.wikimedia.org/wikipedia/de/d/d1/OpenCL_Memory_model.svg
```

Drill down into execution genealogy...

```
1 img("https://upload.wikimedia.org/wikipedia/de/d/d1/OpenCL_Memory_model.svg")
```

• Analyzing stealth implementations

OpenCL Platforms and Devices ↵

- Platforms are OpenGL implementations, listed in /etc/OpenCL/vendors
- Devices are actual CPUs/GPUs
- ICD allows to change platform at runtime

1 [OpenCL.versioninfo\(\)](#)

```
OpenCL.jl version 0.10.9

Toolchain:
- Julia v1.12.5
- OpenCL: 2024.10.24+1
- SPIRV_LLVM_Backend: 20.1.5+3

Julia packages:
- GPUArrays: 11.4.1
- GPUCompiler: 1.8.2
- KernelAbstractions: 0.9.40
- LLVM: 9.4.6
- SPIRVIntrinsics: 0.5.7

Available platforms: 1
- Portable Computing Language
  OpenCL 3.0, PoCL 7.1 Linux, Release, RELOC, SPIR-V, LLVM 20.1.2jl, SLEEF,
  DISTRO, POCL_DEBUG
    . cpu-haswell-AMD EPYC 7763 64-Core Processor (svm:c+f, usm:h+d, bda, fp64,
  il)
```

See also `clinfo` command line tool and `examples/OpenCL/common/device_info.c`.

Tip

► **tl;dr To refresh the list of platforms, you need to quit Julia and open a new session**

Important stats ↗

- Platform
 - name: Portable Computing Language
 - profile: FULL_PROFILE
 - vendor: The pocl project
 - version: PoCL 7.1 Linux, Release, RELOC, SPIR-V, LLVM 20.1.2jl, SLEEF, DISTRO, POCL_DEBUG
- Device
 - name: cpu-haswell-AMD EPYC 7763 64-Core Processor
 - type: cpu

<u>clGetDeviceInfo</u>	Value
CL_DEVICE_GLOBAL_MEM_SIZE	11.62 GiB
CL_DEVICE_MAX_COMPUTE_UNITS	4
CL_DEVICE_LOCAL_MEM_SIZE	512.00 KiB
CL_DEVICE_MAX_WORK_GROUP_SIZE	4096
CL_DEVICE_NATIVE_VECTOR_WIDTH_HALF	0
CL_DEVICE_NATIVE_VECTOR_WIDTH_FLOAT	8
CL_DEVICE_NATIVE_VECTOR_WIDTH_DOUBLE	4
CL_DEVICE_MAX_CLOCK_FREQUENCY	3227 MHz
CL_DEVICE_PROFILING_TIMER_RESOLUTION	1.000 ns

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

Examples ↗

Vectorized sum ↗

```
--kernel void vadd(
    __global const float *a,
    __global const float *b,
    __global float *c,
    int verbose) {
    int i = get_global_id(0);
    c[i] = a[i] + b[i];
```

vadd_size = 512

vadd_verbose = 0

```
> 5 warnings generated.
CL_KERNEL_WORK_GROUP_SIZE           | 4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE   | (0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE            | 0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE          | 0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE | 8
Send command from host to device    | 1.312 µs
Including data transfer             | 40.322 ms
Execution of kernel                 | 58.739 µs
```

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

vadd (generic function with 1 method)

Mandelbrot

```
--kernel void mandelbrot(__global float2 *q,
__global ushort *output, ushort const maxit) {

int gid = get_global_id(0), it;
if (gid == 0)
    printf("%d\n", get_num_groups(0));
float tmp, real = 0, imag = 0;
output[gid] = 0;
for(it = 0; it < maxit; it++) {
    tmp = real * real - imag * imag + q[gid].x;
    imag = 2 * real * imag + q[gid].y;
    real = tmp;
    if (real * real + imag * imag > 4.0f)
        output[gid] = it;
}
```

mandel_size =  512

maxiter =  100

```
1 q = [ComplexF32(r,i) for i=1:-(2.0/mandel_size):-1, r=-1.5:(3.0/mandel_size):0.5];
```

Critical alert

```
ArgumentError: Illegal conversion of a OpenCL.cl.UnifiedDeviceMemory to a  
Ptr{ComplexF32}
```

Execution genealogy

Here's the operational timeline, prioritized by recency:

1. `convert(T::Type{...}, mem::OpenCL.cl.UnifiedDeviceMemory) ...show types...`
sourced from `(memory.jl:8)`
2. `unsafe_convert(P::Type{...}, mem::OpenCL.cl.UnifiedDeviceMemory) ...show types...`
sourced from `(memory.jl:15)`
3. `unsafe_clconvert(typ::Type{...}, mem::OpenCL.cl.UnifiedDeviceMemory) ...show types...`
sourced from `(kernel.jl:440)`
4. macro expansion
sourced from `(kernel.jl:352)`
5. macro expansion
sourced from
6. `convert_arguments(::OpenCL.cl.var"#276#277"..., ::Type{...}, ::OpenCL.CLArray{...},
::OpenCL.CLArray{...}, ::Int64) ...show types...`
sourced from
7. `#clcall#274(::OpenCL.cl.Kernel, ::Type{...}, ::OpenCL.CLArray{...},
::OpenCL.CLArray{...}, ::Int64; kwargs::@Kwargs{}) ...show types...`
sourced from `(kernel.jl:339)`
8. `anonymous function() ...show types...`
sourced from `(Adjacent statement: directive 16)`

```
14 cl.queue!(:profile) do
15     for _ in 1:num_runs
16         evt = clcall(kernel, args...; kws...)
17         wait(evt)
18 
```

cell preview
9. Expand visibility...

```
1 mandel_image = mandel(q, maxiter, mandel_device; global_size=length(q));
```

1 warning generated.	
CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

Critical alert

Another statement defining `mandel_image` exhibits suboptimal performance indicators.

1 aside([CairoMakie.image](#)([CairoMakie.rotr90](#)(`mandel_image`)), `v_offset = -400`)
• Analyzing stealth implementations

mandel (generic function with 1 method)

```
1 function mandel(q::Array{ComplexF32}, maxiter::Int64, device; kws...)
2   cl.device!(device)
3   q = CLArray(q)
4   o = CLArray{Cushort}(undef, size(q))
5
6   prg = cl.Program(; source = mandel\_source.code) |> cl.build!
7   k = cl.Kernel(prg, "mandelbrot")
8
9   timed\_clcall(k, Tuple{Ptr{ComplexF32}, Ptr{Cushort}, Cushort},
10               q, o, maxiter; kws...)
11
12  return Array(o)
13 end
```

1 `mandel_source = code(Example("OpenCL/mandelbrot/mandel.cl"));`

Compute π ↗

Critical alert

```
ArgumentError: Illegal conversion of a OpenCL.cl.UnifiedDeviceMemory to a  
Ptr{Float32}
```

Execution genealogy

Here's the operational timeline, prioritized by recency:

1. `convert(T::Type{...}, mem::OpenCL.cl.UnifiedDeviceMemory)` [...show types...](#)
sourced from `(memory.jl:8)`
2. `unsafe_convert(P::Type{...}, mem::OpenCL.cl.UnifiedDeviceMemory)` [...show types...](#)
sourced from `(memory.jl:15)`
3. `unsafe_clconvert(typ::Type{...}, mem::OpenCL.cl.UnifiedDeviceMemory)` [...show types...](#)
sourced from `(kernel.jl:440)`
4. `macro expansion`
sourced from `(kernel.jl:352)`
5. `macro expansion`
sourced from
6. `convert_arguments(::OpenCL.cl.var"#276#277"{...}, ::Type{...}, ::Int64, ::Float64,
::OpenCL.cl.LocalMem{...}, ::OpenCL.CLArray{...})` [...show types...](#)
sourced from
7. `#clcall#274(::OpenCL.cl.Kernel, ::Type{...}, ::Int64, ::Float64,
::OpenCL.cl.LocalMem{...}, ::OpenCL.CLArray{...}; kwargs::@Kwargs{})` [...show types...](#)
sourced from `(kernel.jl:339)`
8. `anonymous function()` [...show types...](#)
sourced from `(Adjacent statement: directive 16)`

```
14 cl.queue!(:profile) do
15     for _ in 1:num_runs
16         evt = clcall(kernel, args...; kws...)
17         wait(evt)
18     end

```

cell preview

9. Expand visibility...

1 mypi()

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8

► How to compute π with a kernel ?

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

mypi (generic function with 1 method)

First element ↗

Let's write a simple kernel that returns the first element of a vector in global memory.

```
--kernel void first_el(__global float* glob, __global float* result) {
    int item = get_local_id(0);
    if (item == 0)
        *result = glob[item];
}
```

0.9104708f0

1 first_el(rand(Float32, first_el_len))

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8
Send command from host to device	2.675 μ s
Including data transfer	23.116 ms
Execution of kernel	19.005 μ s

first_el (generic function with 1 method)

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

first_el_len = 16

Copy to local memory [🔗](#)

```
--kernel void copy_to_local(__global float* glob, __local float* shared) {  
    int global_size = get_global_size(0);  
    int local_size = get_local_size(0);  
    int item = get_local_id(0);  
    shared[item] = 0;  
    for (int i = 0; i < global_size; i += local_size) {  
        shared[item] += glob[i + item];  
    }  
}
```

1 [copy_to_local\(copy_global_len, copy_local_len\)](#)

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8
Send command from host to device	1.863 µs
Including data transfer	26.215 ms
Execution of kernel	25.829 µs

copy_to_local (generic function with 1 method)

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

copy_global_len = 16

copy_local_len = 16

Reduction on GPU ↵

Many operations can be framed in terms of a MapReduce operation.

- Given a vector of data
- It first map each elements through a given function
- It then reduces the results into a single element

The mapping part is easily embarrassingly parallel but the reduction is harder to parallelize. Let's see how this reduction step can be achieved using arguably the simplest example of `mapreduce`, the sum (corresponding to an identity map and a reduction with `+`).

Sum ↵

```
► (OpenCL = 8.81811, Classical = 8.81811)
1 local_sum(global_len, local_len, local_code, local_device)
```

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8
Send command from host to device	1.202 µs
Including data transfer	48.031 ms
Execution of kernel	15.729 µs

► How to compute the sum an array in local memory with a kernel ?

local_sum (generic function with 1 method)

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

global_len = 16

local_len = 16

Blocked sum ↗

► (OpenCL = 8.81811, Classical = 8.81811)

1 `block_local_sum(block_global_len, block_local_len, factor)`

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8
Send command from host to device	1.523 µs
Including data transfer	35.418 ms
Execution of kernel	25.679 µs

► How to reduce the amount of barrier synchronizations ?

► Was it beneficial in terms of performance for GPUs like in the case of OpenMP ?

block_local_sum (generic function with 1 method)

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

block_global_len = 16

block_local_len = 16

factor = 16

Back to SIMD ↵

- Also called Single Instruction Multiple Threads (SIMT)
- CUDA Warp : width of 32 threads
- AMD waveform : width of 64 threads
- In general : CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE
- Consecutive `get_local_id()` starting from 0
 - So the thread of local id from 0 to 31 are in the same CUDA warp.
- Threads execute the **same instruction** at the same time so no need for barrier .

Warp divergence ↵

Suppose a kernel is executed on a nvidia GPU with `global_size` threads. How much time will it take to execute it ?

```
__kernel void diverge(n)
{
    int item = get_local_id(0);
    if (item < n) {
        do_task_A(); // 'a' ns
    } else {
        do_task_B(); // 'b' ns
    }
}
```




► **How much time will it take to execute it if `global_size` is 32 and `n` is 16 ?**

► **How much time will it take to execute it if `global_size` is 64 and `n` is 32 ?**

Are the threads that are still active in the same warp for you sum example ?

Warp diversion for our sum ↴


► We are still using different warps until the end. Is that a good thing ?



How should we change the sum to keep the working threads on the same warp ?

No warp divergence ↴

Now the same warp is used for all threads so we don't need barrier and it frees other warps to stay idle (reducing power consumption) or do other tasks.



Reordered local sum ↴

```
__kernel void local_sum(__local float* shared)
{
    int items = get_local_size(0);
    int item = get_local_id(0);
    int stride = items / 2;
    float other_val = 0;
    while (stride > 0) {
        barrier(CLK_LOCAL_MEM_FENCE);
        if (item < stride) {
            other_val = 0;
            if (item + stride < items)
                other_val = shared[item+stride];
            shared[item] += other_val;
        }
        stride /= 2;
    }
}
```

► (OpenCL = 8.81811, Classical = 8.81811)

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8
Send command from host to device	531.000 ns
Including data transfer	42.802 ms
Execution of kernel	24.536 µs

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

reordered_global_size = 16

reordered_local_size = 16

SIMT sum ↴

```
__kernel void simt_sum(volatile __local float* shared)
{
    int items = get_local_size(0);
    int item = get_local_id(0);
    barrier(CLK_LOCAL_MEM_FENCE);
    while (items > 1) {
        items /= 2;
        shared[item] += shared[item + items];
    }
}
```

```
► (OpenCL = 2.52133, Classical = 8.81811)
1 local_sum(simt_global_size, simt_local_size, simt_code, simt_device)
```

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8
Send command from host to device	992.000 ns
Including data transfer	30.810 ms
Execution of kernel	23.353 µs

- Why don't we check any condition on item, aren't some thread computing data that won't be used ?

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

simt_global_size =  16

simt_local_size =  16

Beware!

POCL does not synchronize, even for simt_len <= 8

- Why do we need volatile ?

Unrolled sum ↗

- How to get even faster performance by assuming that items is a power of 2 smaller than 512 and that the SIMT width is 32 ?

```
► (OpenCL = 2.52133, Classical = 8.81811)
```

```
1 local_sum(unrolled_global_size, unrolled_local_size, unrolled_code, unrolled_device)
```

CL_KERNEL_WORK_GROUP_SIZE	4096
CL_KERNEL_COMPILE_WORK_GROUP_SIZE	(0, 0, 0)
CL_KERNEL_LOCAL_MEM_SIZE	0 bytes
CL_KERNEL_PRIVATE_MEM_SIZE	0 bytes
CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	8
Send command from host to device	862.000 ns
Including data transfer	34.929 ms
Execution of kernel	23.775 µs

► How to have portable code using unrolling ?

Portable Computing Language ▾

cpu-haswell-AMD EPYC 7763 64-Core Processor ▾

unrolled_global_size = 16

unrolled_local_size = 16

Utils ↗

```
_pretty_time (generic function with 1 method)
1 _pretty_time(x) = BenchmarkTools.prettytime(minimum(x))
```

```

timed_clcall (generic function with 1 method)
1 function timed_clcall(kernel, args...; kws...)
2   info = cl.work_group_info(kernel, cl.device())
3   # See
4   println("CL_KERNEL_WORK_GROUP_SIZE           | ", info.size)
5   println("CL_KERNEL_COMPILE_WORK_GROUP_SIZE    | ", info.compile_size)
6   println("CL_KERNEL_LOCAL_MEM_SIZE            | ",
7         BenchmarkTools.prettymemory(info.local_mem_size))
8   println("CL_KERNEL_PRIVATE_MEM_SIZE          | ",
9         BenchmarkTools.prettymemory(info.private_mem_size))
10  println("CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE | ",
11      info.preferred_size_multiple)

12  # ':profile' sets 'CL_QUEUE_PROFILING_ENABLE' to the command queue
13  queued_submit = Float64[]
14  submit_start = Float64[]
15  start_end = Float64[]
16  cl.queue!(:profile) do
17    for _ in 1:num_runs
18      evt = clcall(kernel, args...; kws...)
19      wait(evt)
20
21      # See
22      https://registry.khronos.org/OpenCL/sdk/3.0/docs/man/html/clGetEventProfilingInfo.html
23      push!(queued_submit, evt.profile_submit - evt.profile_queued)
24      push!(submit_start, evt.profile_start - evt.profile_submit)
25      push!(start_end, evt.profile_end - evt.profile_start)
26    end
27  end
28  println("Send command from host to device | ${_pretty_time(queued_submit)}")
29  println("Including data transfer        | ${_pretty_time(submit_start)}")
30  println("Execution of kernel           | ${_pretty_time(start_end)})")
31 end

```

num_runs =

Activating project at `~/work/LINMA2710/LINMA2710/Lectures`

```
1 using OpenCL, pocl_jll # 'pocl_jll' provides the POCL OpenCL platform for CPU  
devices
```